

Handling Class Imbalance in Random Forest Using Resampling and Cost-Sensitive Learning

Haozhe (Howard) Zeng

Handling Class Imbalance in Random Forest Using Resampling and Cost-Sensitive Learning

Class imbalance is a common challenge in machine learning, particularly in classification tasks. This document demonstrates the use of **resampling techniques** (e.g., oversampling and undersampling) and **cost-sensitive learning methods** to address imbalanced data in random forest models.

Key Topics Covered:

1. **Bagging-Based Methods:**
 - SMOTEBagging, RUSBagging, ROSBagging, Random Balance Bagging (RBBagging)
 2. **Boosting-Based Methods:**
 - SMOTEBoost, RUSBoost, AdaBoost, Cost-Sensitive AdaBoost (AdaC2)
 3. **Specialized Ensemble Methods:**
 - EasyEnsemble, BalanceCascade
 4. **Hybrid Methods:**
 - SMOTETomek (SMOTE combined with Tomek link removal)
-

Required Libraries

Before proceeding, ensure the necessary packages are installed and loaded:

```
# Load dataset
data("PimaIndiansDiabetes")
pima <- PimaIndiansDiabetes

# Prepare features and labels
x <- pima %>% select(-diabetes)
y <- as.factor(ifelse(pima$diabetes == "pos", 1, 0))
table(y) # Check class distribution

## y
##   0   1
## 500 268
```

Bagging-Based Methods (bbagging)

The `bbagging` function implements bagging-based resampling methods, including:

- Random Under-Sampling (RUSBagging)

- Random Over-Sampling (ROSBagging)
- SMOTE (Synthetic Minority Oversampling Technique) Bagging
- Random Balance Bagging (RBBagging)

Parameters for bbaging:

- **x**: A data frame containing the predictor variables.
- **y**: A factor representing the response variable.
- **numBag**: The number of bagging iterations to perform. Default is 10.
- **type**: The type of bagging method to use. Options include:
 - "SMOTEBagging": Uses SMOTE for oversampling.
 - "RUSBagging": Applies random undersampling.
 - "ROSBagging": Performs random oversampling.
 - "RBBagging": Uses random balance bagging.

Example: SMOTEBagging

```
# Train SMOTEBagging model
model <- bbaging(x, y, numBag = 10, type = "SMOTEBagging")

# Predictions
predictions_prob <- predict(model, x, type = "probability")[, 2]
predictions_label <- ifelse(predictions_prob > 0.5, 1, 0)

# Calculate metrics
metrics <- calculate_metrics(y, predictions_label, predictions_prob)
print(metrics)
```

```
## $ConfusionMatrix
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 405  43
##           1  95 225
##
##           Accuracy : 0.8203
##           95% CI : (0.7913, 0.8468)
##           No Information Rate : 0.651
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6216
##
## Mcnemar's Test P-Value : 1.416e-05
##
##           Sensitivity : 0.8100
##           Specificity : 0.8396
##           Pos Pred Value : 0.9040
##           Neg Pred Value : 0.7031
##           Prevalence : 0.6510
##           Detection Rate : 0.5273
##           Detection Prevalence : 0.5833
##           Balanced Accuracy : 0.8248
##
##           'Positive' Class : 0
##
```

```
##
## $Accuracy
## [1] 0.8203125
##
## $WeightedAccuracy
## [1] 0.8035714
##
## $Precision
## [1] 0.8395522
##
## $Recall
## [1] 0.703125
##
## $F1
## [1] 0.7653061
##
## $Specificity
## [1] 0.9040179
##
## $GMean
## [1] 0.7972688
##
## $ROCAUC
## Area under the curve: 0.8874
```

Boosting-Based Methods (bboost)

The `bboost` function applies boosting with resampling or cost-sensitive approaches, such as:

- AdaBoost
- SMOTEBoost
- RUSBoost
- Cost-Sensitive AdaBoost (AdaC2)

Parameters for `bboost`:

- **x**: A data frame containing the predictor variables.
- **y**: A factor representing the response variable.
- **iter**: The number of boosting iterations. Default is 20.
- **type**: The type of boosting method to use. Options include:
 - "AdaBoost": Standard AdaBoost.
 - "SMOTEBoost": Combines boosting with SMOTE.
 - "RUSBoost": Combines boosting with random undersampling.
 - "AdaC2": Cost-sensitive AdaBoost.

Example: SMOTEBoost

```
# Train SMOTEBoost model
model <- bboost(x, y, iter = 20, type = "SMOTEBoost")

# Predictions
predictions_prob <- predict(model, x, type = "probability")[, 2]
predictions_label <- ifelse(predictions_prob > 0.5, 1, 0)

# Calculate metrics
```

```
metrics <- calculate_metrics(y, predictions_label, predictions_prob)
print(metrics)
```

```
## $ConfusionMatrix
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 484  13
##           1  16 255
##
##           Accuracy : 0.9622
##           95% CI : (0.9462, 0.9746)
##           No Information Rate : 0.651
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.9171
##
## Mcnemar's Test P-Value : 0.7103
##
##           Sensitivity : 0.9680
##           Specificity : 0.9515
##           Pos Pred Value : 0.9738
##           Neg Pred Value : 0.9410
##           Prevalence : 0.6510
##           Detection Rate : 0.6302
##           Detection Prevalence : 0.6471
##           Balanced Accuracy : 0.9597
##
##           'Positive' Class : 0
##
##
## $Accuracy
## [1] 0.9622396
##
## $WeightedAccuracy
## [1] 0.9574012
##
## $Precision
## [1] 0.9514925
##
## $Recall
## [1] 0.9409594
##
## $F1
## [1] 0.9461967
##
## $Specificity
## [1] 0.9738431
##
## $GMean
## [1] 0.95726
##
## $ROCAUC
```

```
## Area under the curve: 0.9959
```

EasyEnsemble

EasyEnsemble creates multiple balanced datasets by undersampling the majority class and training individual classifiers.

Parameters for EasyEnsemble:

- **x**: A data frame containing the predictor variables.
- **y**: A factor representing the response variable.
- **iter**: The number of ensemble iterations. Default is 4.
- **allowParallel**: A logical indicating whether to enable parallel computation. Default is **FALSE**.

Example: EasyEnsemble

```
# Train EasyEnsemble model
model <- EasyEnsemble(x, y, iter = 4)

# Predictions
predictions_prob <- predict(model, x, type = "probability")[, 2]
predictions_label <- ifelse(predictions_prob > 0.5, 1, 0)

# Calculate metrics
metrics <- calculate_metrics(y, predictions_label, predictions_prob)
print(metrics)
```

```
## $ConfusionMatrix
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 456    0
##           1  44 268
##
##           Accuracy : 0.9427
##           95% CI : (0.9238, 0.9581)
##       No Information Rate : 0.651
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8785
##
##  McNemar's Test P-Value : 9.022e-11
##
##           Sensitivity : 0.9120
##           Specificity : 1.0000
##       Pos Pred Value : 1.0000
##       Neg Pred Value : 0.8590
##           Prevalence : 0.6510
##       Detection Rate : 0.5938
##  Detection Prevalence : 0.5938
##       Balanced Accuracy : 0.9560
##
##           'Positive' Class : 0
##
```

```
##
## $Accuracy
## [1] 0.9427083
##
## $WeightedAccuracy
## [1] 0.9294872
##
## $Precision
## [1] 1
##
## $Recall
## [1] 0.8589744
##
## $F1
## [1] 0.9241379
##
## $Specificity
## [1] 1
##
## $GMean
## [1] 0.9268087
##
## $ROCAUC
## Area under the curve: 0.9958
```

Balance Cascade

Balance Cascade iteratively trains classifiers while removing easy-to-classify majority instances.

Parameters for BalanceCascade:

- **x**: A data frame containing the predictor variables.
- **y**: A factor representing the response variable.
- **iter**: The number of cascade iterations. Default is 4.

Example: Balance Cascade

```
# Train BalanceCascade model
model <- BalanceCascade(x, y, iter = 4)

# Predictions
predictions_prob <- predict(model, x, type = "probability")[, 2]
predictions_label <- ifelse(predictions_prob > 0.5, 1, 0)

# Calculate metrics
metrics <- calculate_metrics(y, predictions_label, predictions_prob)
print(metrics)
```

```
## $ConfusionMatrix
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 483    0
##           1  17 268
```

```

##
##           Accuracy : 0.9779
##           95% CI   : (0.9648, 0.9871)
##    No Information Rate : 0.651
##    P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.952
##
##    McNemar's Test P-Value : 0.0001042
##
##           Sensitivity : 0.9660
##           Specificity : 1.0000
##    Pos Pred Value : 1.0000
##    Neg Pred Value : 0.9404
##           Prevalence : 0.6510
##    Detection Rate : 0.6289
##    Detection Prevalence : 0.6289
##    Balanced Accuracy : 0.9830
##
##    'Positive' Class : 0
##
##
## $Accuracy
## [1] 0.9778646
##
## $WeightedAccuracy
## [1] 0.9701754
##
## $Precision
## [1] 1
##
## $Recall
## [1] 0.9403509
##
## $F1
## [1] 0.9692586
##
## $Specificity
## [1] 1
##
## $GMean
## [1] 0.9697169
##
## $ROCAUC
## Area under the curve: 1

```

Hybrid Methods: SMOTETomek

SMOTETomek combines SMOTE oversampling with Tomek link removal for better balancing of the dataset.

Parameters for SMOTETomek:

- **x**: A data frame containing the predictor variables.
- **y**: A factor representing the response variable.

- **percOver**: The percentage of oversampling to apply. Default is 100.
- **k**: The number of nearest neighbors to use in SMOTE. Default is 5.

Example: SMOTETomek

```
# Plot original class distribution
print("Before")

## [1] "Before"

table(y)

## y
##   0   1
## 500 268

# Apply SMOTETomek
balanced_data <- SMOTETomek(x, y, percOver = 100)

# Plot new class distribution
print("After")

## [1] "After"

table(balanced_data$y)

##
##   0   1
## 472 508
```